Interrupt is coming to London on October 13! Get tickets

LangChain overview

Copy page

Agent = Model + Harness. LangChain provides `create_agent` : a minimal, highly configurable harness. The harness is everything around the model loop: the prompt, the tools, and any middleware that shapes behavior. Start with the primitives and compose exactly what your use case needs. Supports OpenAI, Anthropic, Google, and more.

LangChain vs. LangGraph vs. Deep Agents

Start with Deep Agents for a "batteries-included" agent with features like automatic context compression, a virtual filesystem, and subagent-spawning. Deep Agents are built on LangChain agents which you can also use directly.

Use LangChain (`create_agent`) for a highly customizable harness, easily tailored to your use case and data.

Use LangGraph, our low-level orchestration framework, for advanced needs combining deterministic and agentic workflows.

Use LangSmith to trace, debug, and evaluate agents built with any of these frameworks. Follow the tracing quickstart to get set up. We recommend you also set up LangSmith Engine which monitors your traces, detects issues, and proposes fixes.

Create an agent

This example demonstrates how to create a simple LangChain agent with a custom tool:

Interrupt is coming to London on October 13! Get tickets


```
    """Get weather for a given city."""
    return f"It's always sunny in {city}!"

agent = create_agent(
    model="openai:gpt-5.5",
    tools=[get_weather],
    system_prompt="You are a helpful assistant",
)

result = agent.invoke(
    {"messages": [{"role": "user", "content": "What's the weather in San Francisc
)
print(result["messages"][-1].content_blocks)
```

See the Installation instructions and Quickstart guide to get started building your own agents and applications with LangChain.

Use LangSmith to trace requests, debug agent behavior, and evaluate outputs. Set `LANGSMITH_TRACING=true` and your API key to get started.

Core benefits

Standard model interface

Use one interface for chat models, embeddings, and more across providers. Switch models with minimal code changes and keep your

Highly configurable harness

Start with `create_agent` as a minimal harness and add capabilities incrementally through middleware. Compose only what your use case

Interrupt is coming to London on October 13! Get tickets

LangChain Docs

Built on top of LangGraph

LangChain's agents are built on top of LangGraph. This allows us to take advantage of LangGraph's durable execution, human-in-the-loop support, persistence, and more.

Learn more >

Debug with LangSmith

Inspect traces, tool calls, state transitions, and latency in one place. Find failure modes, evaluate quality, and improve agent behavior with execution data.

Learn more >

Connect these docs to your agent of choice via MCP for real-time answers.

Edit this page on GitHub or file an issue.

Was this page helpful? Yes No

LangChain Docs

Resources

Company

Interrupt is coming to London on October 13! Get tickets

LangChain Docs